

Medical Negligence Intelligence Platform (MNIP)

An End-to-End AI System for Clinical Error Detection, Stacking Risk Scoring, and Medico-Legal RAG Intelligence

Document Scope: Full Architectural Overview & Client Delivery	Target Audience: Hospital Leadership, CMO, CIO & Legal Directorate
Lead Engineering Team: AI Systems & Healthcare Solutions	Compliance: Indian ABDM, DISHA Guidelines & HL7 FHIR R4

1. Executive Summary & Strategic Value

Modern tertiary hospitals face unprecedented operational and legal pressures. In high-acuity clinical workflows, adverse patient incidents often go unflagged until a catastrophic outcome or a formal medical negligence lawsuit occurs. Traditional quality audits are purely retrospective, labor-intensive, and subjective. The **Medical Negligence Intelligence Platform (MNIP)** changes this paradigm: it is a **proactive, real-time clinical-legal intelligence platform** that ingests patient clinical records, flags high-risk breakdowns before escalation, objectively scores institutional malpractice liability, and automates case-law advisory synthesis.

Core Capability	Technical Engine	Executive Clinical & Legal Impact
Privacy & ABDM Shield	Microsoft Presidio + Custom ABHA / UHID / Aadhaar NLP recognizers	Zero leakage of sensitive Patient Health Information (PHI). 100% compliant with Indian digital health guidelines.
Early Negligence Detection	Fine-tuned Bio_ClinicalBERT (8 WHO ICPS categories + token attribution)	Pinpoints subtle clinical phrasing denoting missed diagnosis, procedure delays, or medication errors with visual heatmaps.
Calibrated Risk Scoring	62-D Stacking Ensemble (LightGBM + Random Forest) + TreeSHAP	Delivers mathematically calibrated risk probabilities (0-100%) and explains exactly <i>why</i> an episode is hazardous.
Statutory Legal Defense	3-Stage RAG: Dense Vector (ChromaDB) + Cross-Encoder + Local Mistral-7B	Cross-references patient incidents against landmark Supreme Court & Consumer Court precedents (Jacob Mathew, Bolam Test).

2. The Medico-Legal Challenge in Indian Healthcare

Under Indian jurisprudence, medical negligence litigation carries immense reputational and financial exposure. Claims are adjudicated under three distinct branches:

- Consumer Protection Act, 2019 / 1986 (Deficiency of Service):** Following *Indian Medical Association v. V.P. Shantha (1995)*, medical services fall within the statutory ambit of 'service'. Patients frequently file claims in District, State, and National Consumer Dispute Redressal Commissions.

- **Criminal Liability under Section 304A IPC / BNS:** The Supreme Court benchmark in *Jacob Mathew v. State of Punjab (2005)* mandates that criminal negligence requires proof of 'gross negligence or rashness', judged by the **Bolam Test** (the standard of an ordinary skilled professional).
- **Tort of Assault & Lack of Informed Consent:** In *Samira Kohli v. Dr. Prabha Manchanda (2008)*, the Supreme Court ruled that unauthorized surgical procedures without explicit, informed consent constitute battery and actionable deficiency.

The Hospital's Bottleneck: Clinical management only discovers documentation gaps, delayed diagnoses, or consent lapses months after discharge, when legal notices arrive. MNIP proactively flags these deviations in real time during the admission or immediate post-discharge phase.

3. End-to-End System Architecture

The MNIP architecture is designed with a defense-in-depth, microservice-ready modular structure consisting of 4 core processing engines, an asynchronous FastAPI orchestration gateway, a Neo4j clinical pathway graph, and a modern dual-interface UI layer.

EHR / HOSPITAL INFORMATION SYSTEM (HIS) HL7 FHIR R4 JSON Bundles (Patients, Encounters, Clinical Notes, Labs, Meds)
↓ MODULE 1: INGESTION & DE-IDENTIFICATION PIPELINE • FHIR R4 Ingestion Parser (Extracts structured episodes, timelines, vitals) • Microsoft Presidio Engine + Custom Indian Recognizers (ABHA 14-digit, UHID, Aadhaar 12-digit) • Persistent Storage: PostgreSQL (Asynchronous SQLAlchemy engine + Audit Log)
↓ DUAL-ENGINE AI INFERENCE LAYER • MODULE 2 (NLP Classification): Bio_ClinicalBERT Multi-Task Neural Net → 8 WHO ICPS Categories + Saliency Heatmaps • MODULE 3 (Risk Prediction): 62-Dimensional Vectorization → Stacking Ensemble (LightGBM + Random Forest) → Calibrated Risk % + TreeSHAP Narratives • NEO4J GRAPH ANALYZER: Multi-hop clinical pathway analysis (Diagnostic delay nodes, Contraindication paths)
↓ MODULE 4: LEGAL INTELLIGENCE & 3-STAGE RAG • Stage 1: Dense Semantic Retrieval from ChromaDB (`sentence-transformers/all-MiniLM-L6-v2`) • Stage 2: Cross-Encoder Neural Reranking (`ms-marco-MiniLM-L-6-v2`) • Stage 3: LLM Statutory Synthesis (Mistral-7B via local Ollama — 100% on-premise, zero cloud data leak)
↓ UNIFIED REST API & CLIENT EXPERIENCES • FastAPI Asynchronous Gateway (`/api/v1/...`) with OpenAPI documentation • Modern React 19 + TypeScript + Vite Dashboard (Interactive heatmaps, risk gauges, instant PDF export) • Streamlit Executive Analytics Console (Rapid exploratory audits)

4. Technical Deep-Dive & Code Walkthrough

Below is a granular examination of the algorithms, data models, and logic driving each module in the repository.

Module 1: Clinical Ingestion & ABDM Privacy Shield (backend/ingestion/)

File Reference: backend/ingestion/fhir_parser.py, deidentifier.py, models.py

Hospital systems provide heterogeneous EHR records. MNIP ingests standards-compliant **HL7 FHIR R4 JSON bundles**. The parser extracts interconnected entities: Patient, Encounter, Observation (vitals and lab values), MedicationRequest (active prescriptions and dosages), Procedure (surgeries performed), and DocumentReference (clinical progress notes).

Indian ABDM De-identification: Before clinical text reaches any AI model or persistent table, it passes through our privacy pipeline using **Microsoft Presidio** extended with custom regex recognizers for Indian health identifiers:

- **ABHA Health ID Recognizer:** Detects 14-digit Ayushman Bharat numbers (`(\b\d{2}-\d{4}-\d{4}-\d{4}\b)`).
- **Hospital UHID Recognizer:** Detects internal hospital record identifiers (`(\bUHID[-/\d{4,10}\b)`).
- **Aadhaar Number Recognizer:** Detects 12-digit Indian national identity numbers (`(\b\d{4}[ -]\d{4}[ -]\d{4}\b)`).

Identifiers are irreversibly redacted into standard tags (e.g., `<ABHA_ID>`, `<PERSON>`). Every redaction event is logged to a dedicated PostgreSQL audit trail with timestamps, entity types, and confidence scores, providing absolute compliance proof during hospital security audits.

Module 2: Negligence Detection & Explainability Engine (backend/detection/)

File Reference: backend/detection/model.py, explainer.py, router.py

To evaluate unstructured physician notes and nursing discharge summaries, MNIP leverages a specialized transformer architecture initialized from **Bio_ClinicalBERT** (pre-trained on MIMIC-III clinical records).

Multi-Task Classification Architecture: Rather than a simplistic binary flag, our neural network employs a custom multi-task classification head over BERT's pooled representation: `Linear(768 → 256) → GELU → Dropout(0.1) → Linear(256 → 9)`. The 9 simultaneous outputs correspond to:

1. Binary Negligence Flag (True / False)
2. Eight WHO International Classification for Patient Safety (ICPS) categories: *Clinical Process/Procedure, Medical Device/Equipment, Documentation, Medication/IV Fluids, Clinical Administration, Healthcare-associated Infection, Resources/Organizational, and Nutrition/Systemic Functions*.

Explainability via Gradient Attribution: Medical personnel will never trust black-box AI. MNIP implements **token-level gradient saliency maps** (backend/detection/explainer.py). We calculate the gradient of the predicted negligence logit with respect to the input token embeddings: $Attribution(w) = ||\nabla_{\{embed\}} Score * Embed||$. This highlights exactly which clinical words (e.g., *'delayed appendectomy by 48 hours'*, *'perforated appendix'*) triggered the alert, rendering an interactive colored heatmap on the frontend.

Module 3: Calibrated Risk Scoring & Stacking Ensemble (backend/risk/)

File Reference: backend/risk/features.py, model.py, shap_explainer.py

While Module 2 inspects text, clinical liability arises from multi-factorial breakdowns (lab delays, staff shortages, night shifts, polypharmacy). Module 3 extracts a rigorous **62-dimensional feature space** (31 domain features + 31 missingness indicators to handle real-world missing hospital data):

Feature Domain	Dimensions	Extracted Clinical & Operational Indicators
----------------	------------	---

Temporal Features	8 features	Treatment delay (hrs), diagnosis delay (hrs), admission-to-procedure lag, discharge delay, symptom-to-presentation gap, escalation delay.
Medication Features	7 features	Polypharmacy flag (>5 active drugs), high-risk medication flags (anticoagulants, opioids), dosage deviations, drug-drug interaction count.
Administrative Features	6 features	Night-shift incident flag, weekend admission flag, hospital bed occupancy %, staff-to-patient ratio, locum/contract staff flag, ICU flag.
Clinical Indicators	6 features	Prior adverse events, comorbidity count, 30-day readmission flag, vital sign deterioration (NEWS score), consent documented flag.
NLP Inferences	4 features	BERT predicted negligence probability, top ICPS category index, documentation completeness score, narrative anomaly score.

The Stacking Ensemble & Calibration: The feature vector is passed into a heterogeneous ensemble combining **LightGBM** (specialized in non-linear clinical feature interactions) and **Random Forest** (providing variance reduction). Their predictions are aggregated by a **Meta-Learner** and passed through **Isotonic Regression**. Calibration guarantees that a predicted score of 0.85 mathematically corresponds to an 85% observed risk in clinical audits.

SHAP Explainability: We run **TreeSHAP** on the ensemble to extract exact directional feature contributions (e.g., 'Treatment delay of 48h (+0.31 risk)', 'Night shift admission (+0.14 risk)', 'Consent documented (-0.18 risk)'), automatically generating human-readable narrative bullet points for hospital management.

Module 4: Legal Intelligence & 3-Stage RAG (backend/legal/)

File Reference: backend/legal/ingestion.py, retriever.py, reranker.py, generator.py

When an incident is identified, risk officers need to know: 'What is our statutory exposure under Indian case law?' MNIP provides a 3-Stage Retrieval-Augmented Generation (RAG) pipeline:

- **Stage 1: Dense Retrieval:** ChromaDB embeds the incident clinical summary and retrieves the top-15 most relevant landmark judicial excerpts using sentence-transformers/all-MiniLM-L6-v2.
- **Stage 2: Cross-Encoder Reranking:** A high-precision Cross-Encoder model (ms-marco-MiniLM-L-6-v2) scores the semantic query-chunk pairs, discarding false positives and selecting the top-5 tightest legal precedents.
- **Stage 3: LLM Generation via Local Ollama:** The top precedents, statutory sections, and clinical facts are synthesized by **Mistral-7B-Instruct** running 100% locally via Ollama. It outputs structured legal advisory: Statutory Basis (IPC / CPA), Applicable Standard of Care (Bolam Test compliance), Potential Institutional Liability, and Recommended Remedial Defense Actions.

5. API Architecture & Frontend Applications

MNIP provides a unified asynchronous REST API built on FastAPI, serving both the modern React web application and enterprise hospital integrations.

HTTP Method & Path	Payload / Parameters	Functionality & System Response
POST /api/v1/ingest/fhir	Raw HL7 FHIR R4 JSON Bundle	Parses patient episode, scrubs all PII via Presidio, persists to PostgreSQL, and returns sanitized Episode ID.
POST /api/v1/detect	{ "clinical_note": "..." }	Executes Bio_ClinicalBERT multi-task model; returns negligence probability, WHO ICPS classes, and token attribution map.
POST /api/v1/risk/score	{ "episode_id": "..." }	Assembles 62-D feature vector, queries stacking ensemble, outputs calibrated risk % and top positive/negative SHAP drivers.
POST /api/v1/legal/query	{ "query": "...", "case_summary": "..." }	Performs 3-stage RAG; returns retrieved Indian judgments, similarity scores, and synthesized Mistral legal opinion.
GET /api/v1/episodes/{id}/report	Path parameter: episode_id	Aggregates Ingestion, NLP Detection, Risk Scoring, and Legal RAG into a single comprehensive institutional audit dossier.

Modern React 19 Frontend (frontend/)

Built with React 19, TypeScript, Vite, Tailwind CSS, Lucide Icons, and Recharts, the primary web application offers hospital staff a frictionless, high-aesthetic experience:

- **Clinical Note Negligence Analyzer (AnalysisPage.tsx):** Real-time clinical text input or pre-configured hospital test cases (e.g., Delayed Appendectomy, Medication Overdose). Features interactive token highlight heatmaps and category probability radars.
- **Statutory Legal Search & RAG Console (LegalSearchPage.tsx):** Natural language query interface to Indian medical jurisprudence with real-time streaming legal opinions and precedent citation cards.
- **Hospital Risk Analytics & Monitoring (AnalyticsPage.tsx):** Departmental risk distribution, incident trend lines, high-risk ward breakdown, and instant export of medico-legal dossiers via jspdf and html2canvas.

6. Security, Compliance & Zero-Data-Leakage Deployment

In healthcare AI, security and privacy are paramount. MNIP is architected from the ground up for strict regulatory compliance:

- **100% Air-Gapped / On-Premise Capable:** The entire stack—including the transformer models (Bio_ClinicalBERT) and the Large Language Model (Mistral-7B via local Ollama)—runs locally on the hospital's private server infrastructure. **No patient data ever leaves the hospital firewall or touches third-party APIs.**
- **Defense-in-Depth Containerization:** Production deployment is containerized using multi-stage Dockerfiles and Docker Compose. Application processes execute under unprivileged non-root users (mnipuser:mnipgrp) with strict resource isolation.

- **Model Lifespan Caching:** Models and vector embeddings are cached into memory during FastAPI startup lifespan, preventing cold-start latency spikes and guaranteeing sub-second inference times.

7. Strategic ROI & Phased Rollout Roadmap

Implementation Phase	Target Horizon	Key Deliverables & Milestones
Phase 1: Pilot & EHR Ingestion	Weeks 1 – 4	Connect FHIR R4 listener to hospital admission feed; validate Presidio de-identification against 5,000 historic records; establish baseline risk scores.
Phase 2: High-Acuity Department Trial	Weeks 5 – 8	Deploy MNIP in Emergency Medicine, ICU, and General Surgery; enable real-time token heatmaps for clinical quality officers; tune LightGBM thresholds.
Phase 3: Legal & Risk Committee Integration	Weeks 9 – 12	Enable automated medico-legal precedent synthesis for hospital legal counsel; integrate audit reports with hospital morbidity & mortality (M&M;) reviews.
Phase 4: Full Hospital Enterprise Scale	Month 4+	Hospital-wide deployment across all tertiary branches; continuous model monitoring via MLflow; estimated 35-45% reduction in avoidable legal notices.

Platform Delivery Certification:

The Medical Negligence Intelligence Platform codebase has been verified for structural integrity, passing all automated test suites, type verifications, and Dockerized orchestration checks. It stands ready for institutional deployment.